

# EC349 Project

Wai Yan

2023-12-03

## Tabula statement

We're part of an academic community at Warwick.

Whether studying, teaching, or researching, we're all taking part in an expert conversation which must meet standards of academic integrity. When we all meet these standards, we can take pride in our own academic achievements, as individuals and as an academic community.

Academic integrity means committing to honesty in academic work, giving credit where we've used others' ideas and being proud of our own achievements.

In submitting my work I confirm that:

1. I have read the guidance on academic integrity provided in the Student Handbook and understand the University regulations in relation to Academic Integrity. I am aware of the potential consequences of Academic Misconduct.
2. I declare that the work is all my own, except where I have stated otherwise.
3. No substantial part(s) of the work submitted here has also been submitted by me in other credit bearing assessments courses of study (other than in certain cases of a resubmission of a piece of work), and I acknowledge that if this has been done this may lead to an appropriate sanction.
4. Where a generative Artificial Intelligence such as ChatGPT has been used I confirm I have abided by both the University guidance and specific requirements as set out in the Student Handbook and the Assessment brief. I have clearly acknowledged the use of any generative Artificial Intelligence in my submission, my reasoning for using it and which generative AI (or AIs) I have used. Except where indicated the work is otherwise entirely my own.
5. I understand that should this piece of work raise concerns requiring investigation in relation to any of points above, it is possible that other work I have submitted for assessment will be checked, even if marks (provisional or confirmed) have been published.
6. Where a proof-reader, paid or unpaid was used, I confirm that the proofreader was made aware of and has complied with the University's proofreading policy.
7. I consent that my work may be submitted to Turnitin or other analytical technology. I understand the use of this service (or similar), along with other methods of maintaining the integrity of the academic process, will help the University uphold academic standards and assessment fairness.

## 1. Introduction

In the early 2000s, the significance of online user reviews and their impact on product sales gained widespread recognition (Chevalier and Mayzlin 2006). Leveraging this information source has become instrumental for businesses in understanding consumer behaviour and enhancing profitability. Notably, the application

of machine learning, particularly deep learning, has become prevalent in analysing user reviews, with a particular emphasis on the hospitality and tourism industry (Dickinger and Mazanec 2015; Alaei, Becken, and Stantic 2019). One of the common approaches is sentiment analysis which involves categorising the textual content into positive, neutral, and negative sentiments (Schmunk et al. 2013). To better predict the user reviews, sentiment, captured by the count of positive and negative words in the review text, is introduced as one of the variables, alongside other variables available in the datasets.

## 2. Related work

Literature in review rating prediction using Yelp datasets was well-documented, with an extensive evaluation of both machine learning and deep learning models (Elkouri 2015; Asghar 2016; A. Rafay, M. Suleman, and A. Alim 2020a; S. Liu 2020; Z. Liu 2020). Surprisingly, most researchers found that logistic regression models tend to outperform other supervised learning algorithms such as random forest, support vector machine (SVM) and Naive Bayes Classifier. They yield accuracy ranged between 64-65% with some variations in training times (Elkouri 2015; S. Liu 2020; Z. Liu 2020; Asghar 2016). However, there are mixed findings in the comparison between machine learning and deep learning methods<sup>1</sup>.

Additionally, similar projects employed F1 score as an evaluation metric, accounting for Type I and Type II errors, along with confusion matrix (S. Liu 2020; Z. Liu 2020; Asghar 2016). Also, in terms of text pre-processing, the td-idf transformer or vectoriser was used to generate weight of each word, as the length of the text may affect the count for “buzzwords” (Elkouri 2015; Asghar 2016; A. Rafay, M. Suleman, and A. Alim 2020b; S. Liu 2020; Z. Liu 2020). Regrettably, these are not implemented in this data project, as it would be computationally challenging and beyond the learning scope.

## 3. Data setup

### 3.1 Data preparation

In terms of sentiment analysis, a structured pre-processing approach, as suggested in previous literature, was employed to enhance the classification of textual data (Noori 2021). Using a curated list of prevalent positive and negative words commonly employed by online users<sup>2</sup>, the text in review data was converted into lower case, with punctuation, digits and English “Stop words”<sup>3</sup> removed. Here are some examples of negative and positive words:

Positive words	Negative words
winning	rankle
exceedingly	slowed
magnanimous	pitiless
unforgettable	skeptically
faithfulness	unreasonably
sweet	stubbornly
amiably	scandalized
agreeably	mendacious
merciful	disrupt
ethical	imbecile

<sup>1</sup>Mostly Long Short Term Memory (LSTM) and Bidirectional Encoder Representations from Transformers (BERT)

<sup>2</sup>Available online at github:


- positive-text.txt: <https://gist.github.com/mkulakowski2/4289437#file-positive-words-txt-L9>

- negative-text.txt: <https://gist.github.com/mkulakowski2/4289441>

<sup>3</sup>Words like ‘the’, ‘that’, ‘is’ etc, occur frequently across reviews and are not very useful.

### 3.2 Feature selection

To identify relevant variables for analysis, preliminary multinomial logistic regressions were conducted, and graphs were generated to assess the associations between potential predictors and the outcome variable, which ranges from 1 to 5. The bar graphs below show the distribution of votes: useful, cool and funny.



Initially, variables such as word count of text, check-in count, sum of compliments of each user, and review count (both user and business) were considered. However, subsequent analysis revealed insignificant coefficients, leading to their exclusion from the final model.

### 3.3 Missing values

As the user, review and business data are merged, the missing values are omitted and the types of some variables are changed to better conduct analysis and interpret results. The training data now consists of 269878 observations (originally 1388056). The final list of variables are as below:

```
'data.frame':  279878 obs. of  9 variables:
 $ stars.x      : Factor w/ 5 levels "1","2","3","4",...: 2 1 4 2 4 1 4 1 5 4 ...
 $ useful       : num  0 10 1 3 0 7 0 1 0 3 ...
 $ funny        : num  1 11 0 0 0 6 0 0 1 1 ...
 $ cool         : num  0 2 1 0 0 5 0 0 0 4 ...
 $ positive_word_count: num  5 10 6 2 2 8 3 1 2 3 ...
 $ negative_word_count: num  3 22 2 3 4 5 1 2 1 2 ...
 $ sentiment    : Factor w/ 3 levels "negative","neutral",...: 3 1 3 1 1 3 3 1 3 3 ...
 $ average_stars : num  4.18 3.83 4.31 2.75 4.5 4.07 3.41 4.04 3.69 4.07 ...
 $ stars.y      : num  3.5 2 4.5 3 3 4 4 4 3.5 4.5 ...
```

where stars.x is the user rating and stars.y is the business rating.

## 4. Methodology

### 4.1 Identification strategy: Random Forest

Random forest provides a comprehensive prediction model to prevent overfitting due to sensitivity towards outliers and noises (Breiman 2001). Imbalanced class, characterised by huge differences in class probabilities,

can impact the performance of prediction model especially decision trees (Japkowicz and Stephen 2002; Chawla 2003; Muchlinski et al. 2016). In this case where the ratings of 4 and 5 are exceptionally high compared to other ratings (as shown below), the use of random forest can be rationalised.

Table 2: Actual frequency

Actual_Outcome	Actual_Frequency
1	1547
2	803
3	1012
4	2052
5	4586

The internal feature selection and decorrelation of random forest are also applicable in this case where predictors are found to be higher correlated, indicating a high possibility of multicollinearity<sup>4</sup>. This is because random forest strategically limits predictor consideration at each split, mitigating the impact of a single dominant predictor (James et al. 2021).

## 4.2 Application

The model is specified below:

- Outcome: stars.x (user review)
- Predictors: useful, funny, cool, positive\_word\_count, negative\_word\_count, sentiment, average\_stars, stars.y (business review)
- Number of trees<sup>5</sup>: 100
- The number of variables at each split, m is set to be the square root of number of predictors, as suggested in most general cases
- The node size is set to be 10

## 4.3 Results and model evaluation

The confusion matrix which shows the the number of correct and false predictions as well as accuracy are computed to evaluate the performance of the model. The model managed to produce 59.66% of correct prediction. The confusion matrix presented below summarises the performance of the random forest model:

Table 3: Random Forest Confusion Matrix

Prediction	X1.star	X2.stars	X3.stars	X4.stars	X5.stars
NA	1	2	3	4	5
1	1242	355	174	85	91
2	87	108	47	42	31
3	46	58	96	71	41
4	54	113	273	473	376
5	118	169	422	1381	4047

<sup>4</sup>Variance Inflation Factor (VIF) is used to detect multicollinearity between predictors, values as high as 27 were obtained.

<sup>5</sup>Initial attempt was 500 trees but this was not allowed due to limited memory capacity.

Despite achieving a marginally higher accuracy, the model’s limitations in predicting ratings between 2 and 3 persisted, with results skewed towards ratings 1 and 5. Additionally, the constraint on the number of trees may have influenced the model’s predictive performance, especially considering the scale of the dataset.

The out-of-bag estimation error, standing at 40.54%, raises concerns about the generalisation performance of the random forest model. This substantial error suggests a potential risk of overfitting, where the model could be influenced by noise.

#### 4.4 Comparison with Logistic Regression

Due to the limited interpretability of the random forest, the analysis was replicated using multinomial logistic regression. In this comparison, logistic regression exhibited a slightly diminished performance compared to the random forest, achieving an overall accuracy of 58.65%. These findings align with recent observations in the medical field, where no significant differences were noted between the performance of logistic regression models and random trees(Christodoulou et al. 2019).

In the logistic regression analysis, the advantage lies in its ability to show the interaction between the outcome variable and predictors. Most predictors appeared to be able to identify underlying patterns within the review data, with a certain level of ambiguity in distinguishing between 4 and 5-star ratings.

Table 4: Logistic Regression Coefficients

X.Intercept	useful	funny	cool	positive_word	negative_word	sentiment	neutral	positive	average_stars	y
-	-	-	0.2498174	0.1401503	-0.1105936	0.2365958	0.4106026	0.6643919	0.5061668	
4.577626	0.0832350	0.0422358								
-	-	-	0.7035952	0.2566341	-0.2964977	0.3281336	0.7015634	1.2324862	0.7499524	
7.397184	0.2337800	0.1580130								
-	-	-	0.9819081	0.3353011	-0.5305205	0.4222237	1.2211666	1.6748479	1.1271462	
10.191325	0.3337514	0.2964029								
-	-	-	0.9673846	0.3423598	-0.7066851	0.2542999	1.2775829	2.4241066	1.8670458	
15.139453	0.3306010	0.3192753								

Table 5: Logistic Regression Confusion Matrix

Prediction	X1.star	X2.stars	X3.stars	X4.stars	X5.stars
NA	1	2	3	4	5
1	1298	433	192	91	104
2	34	39	31	24	12
3	20	27	53	33	23
4	66	85	208	281	253
5	129	219	528	1623	4194

Despite being slightly better in predicting polarity with ratings of 1 and 5, it tends to present bias towards dominant classes, resulting in less accurate predictions for 2 and 3-star ratings compared to random forest, justifying the final choice of random forest model.

## 5. Discussion and Conclusion

The data project faces challenges with prolonged training times and high memory requirements for both random trees and logistic regression, highlighting the need for simplified models that balance processing

efficiency with performance. Additionally, it highlighted the need for a standardised approach to address imbalanced data as some existing suggestions like re-sampling and improved weights may lead to missing information and increased computational difficulties (Japkowicz and Stephen 2002; Shahhosseini and Hu 2020).

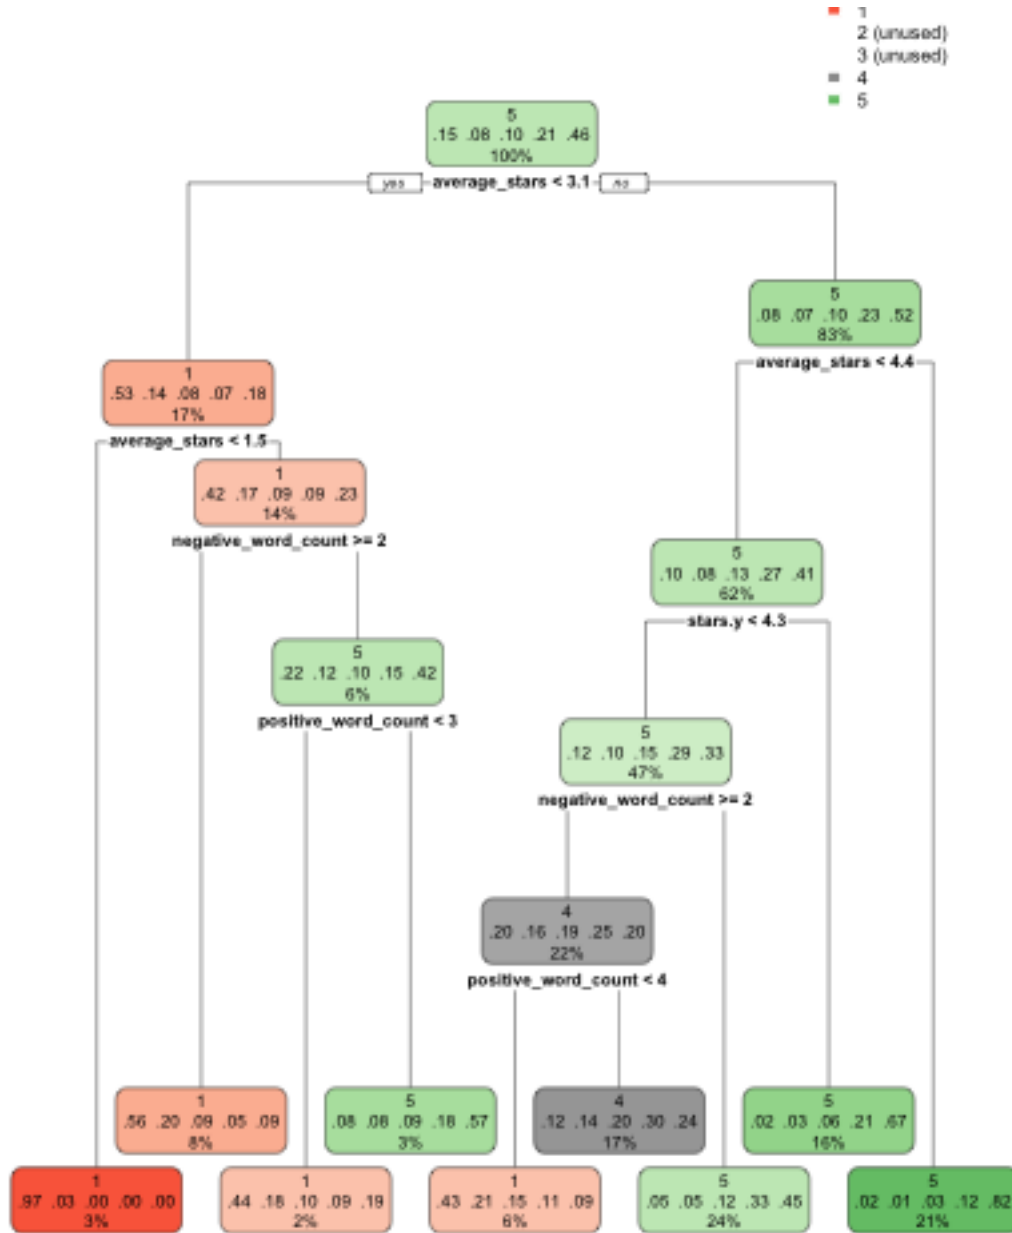
In conclusion, the analysis indicates that random forest provided a higher accuracy in classifying imbalanced classes compared to logistic regression, contrary to previous literature. Despite this advantage, the random forest model does exhibit some misclassification errors, which may be attributed to overfitting. However, its distinct benefits, including enhanced generalisation and decorrelation, underscore its suitability for this specific application. Considering overall accuracy and potential overfitting, random choice remains the most effective model to the best of my knowledge.

(Word count: 1124)

## References

- A. Rafay, M. Suleman, and A. Alim. 2020b. “2020 International Conference on Emerging Trends in Smart Technologies (ICETST).” In, 8138–43. <https://doi.org/10.1109/ICETST49965.2020.9080713>.
- . 2020a. “2020 International Conference on Emerging Trends in Smart Technologies (ICETST).” In, 8138–43. <https://doi.org/10.1109/ICETST49965.2020.9080713>.
- Alaei, Ali Reza, Susanne Becken, and Bela Stantic. 2019. “Sentiment Analysis in Tourism: Capitalizing on Big Data.” *Journal of Travel Research* 58 (2): 175–91. <https://doi.org/10.1177/0047287517747753>.
- Asghar, Nabiha. 2016. “Yelp Dataset Challenge: Review Rating Prediction.” *arXiv.org*, May. <http://0-search.proquest.com.pugwash.lib.warwick.ac.uk/working-papers/yelp-dataset-challenge-review-rating-prediction/docview/2079845554/se-2?accountid=14888>.
- Breiman, Leo. 2001. “Random Forests.” *Machine Learning* 45 (1): 5–32. <https://doi.org/10.1023/A:1010933404324>.
- Chawla, Nitesh V. 2003. “C4. 5 and Imbalanced Data Sets: Investigating the Effect of Sampling Method, Probabilistic Estimate, and Decision Tree Structure.” In *Proceedings of the ICML*, 3:66. Citeseer.
- Chevalier, Judith A., and Dina Mayzlin. 2006. “The Effect of Word of Mouth on Sales: Online Book Reviews.” *Journal of Marketing Research* 43 (3): 345–54. <https://doi.org/10.1509/jmkr.43.3.345>.
- Christodoulou, Evangelia, Jie Ma, Gary S. Collins, Ewout W. Steyerberg, Jan Y. Verbakel, and Ben Van Calster. 2019. “A Systematic Review Shows No Performance Benefit of Machine Learning over Logistic Regression for Clinical Prediction Models.” *Journal of Clinical Epidemiology* 110: 12–22. <https://doi.org/https://doi.org/10.1016/j.jclinepi.2019.02.004>.
- Dickinger, Astrid, and Josef A. Mazanec. 2015. “Significant Word Items in Hotel Guest Reviews: A Feature Extraction Approach.” *Tourism Recreation Research* 40 (3): 353–63. <https://doi.org/10.1080/02508281.2015.1079964>.
- Elkouri, Andrew. 2015. “Predicting the Sentiment Polarity and Rating of Yelp Reviews.”
- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2021. *An Introduction to Statistical Learning: With Applications in r*. Second. New York: Springer.
- Japkowicz, Nathalie, and Shaju Stephen. 2002. “The Class Imbalance Problem: A Systematic Study.” *Intelligent Data Analysis* 6 (5): 429–49. <https://doi.org/10.3233/IDA-2002-6504>.
- Liu, Siqi. 2020. “Sentiment Analysis of Yelp Reviews: A Comparison of Techniques and Models.”
- Liu, Zefang. 2020. “Yelp Review Rating Prediction: Machine Learning and Deep Learning Models.” *arXiv.org*, December. <http://0-search.proquest.com.pugwash.lib.warwick.ac.uk/working-papers/yelp-review-rating-prediction-machine-learning/docview/2470136753/se-2?accountid=14888>.
- Muchlinski, David, David Siroky, Jingrui He, and Matthew Kocher. 2016. “Comparing Random Forest with Logistic Regression for Predicting Class-Imbalanced Civil War Onset Data.” *Political Analysis* 24 (1): 87–103.
- Noori, Behrooz. 2021. “Classification of Customer Reviews Using Machine Learning Algorithms.” *Applied Artificial Intelligence* 35 (8): 567–88. <https://doi.org/10.1080/08839514.2021.1922843>.
- Schmunk, Sergej, Wolfram Höpken, Matthias Fuchs, and Maria Lexhagen. 2013. “Sentiment Analysis: Extracting Decision-Relevant Knowledge from UGC.” In, edited by Zheng Xiang and Iis Tussyadiah, 253265. Cham: Springer International Publishing.

## Appendix



Decision tree was also plotted but ratings of 2 and 3 were completely neglected in this model